

# CRITERION

*Evaluate. Compare. Diagnose.*

4

Grader types

1 API

Agent-native via MCP

SQLite

Full eval history

- EXACT MATCH
- CONTAINS MATCH
- REGEX
- LLM JUDGE
- MCP SERVER
- LEADERBOARD

Open-source LLM evaluation platform. POST a prediction, get a score, pass/fail verdict, and LLM reasoning back instantly.

`critterion-evall.streamlit.app`

- EXACT MATCH
- CONTAINS MATCH
- LLM JUDGE
- MCP SERVER

# LLM outputs fail silently in production.

No grading. No history. No patterns. Just a bad response and a guess. Teams need automated, explainable scoring they can query programmatically.

## HOW IT WORKS

**Send a prediction. Get a verdict.**

POST /jobs with your input, prediction, and reference. Get back a score, pass/fail, and LLM reasoning instantly.

## 4 GRADERS

**From exact match to LLM judge.**

Exact match, contains match, regex, and Gemini-powered LLM judge with reasoning output.

## PERSISTENCE

**Every result is stored.**

SQLite backed. Full history via GET /history. Leaderboard ranks models by avg score and pass rate.

## AGENT-NATIVE

**MCP server included.**

Claude and other agents call run\_eval and get\_leaderboard as tools directly. No browser needed.

# 54/55 tests passing on the live API.

98.2%

Pass rate on live Render

50

Concurrent requests, zero corruption

100

Job batch in 0.59s

1MB

Payload handled without crash

- UNICODE + EMOJI

Accented chars and special symbols handled correctly
- SQLITE CONCURRENCY

No lock errors under 50 concurrent writes
- LEADERBOARD MATH

3 out of 4 pass equals exactly 75.0% pass rate
- IDEMPOTENCY

Same eval twice returns identical scores
- PROMPT INJECTION

LLM judge resistant to payload attacks
- MALFORMED INPUT

Missing fields return 422, not 500
- REGEX SAFETY

Invalid patterns return clean error, not 500
- CLUSTERING BOUNDARY

Single text and n greater than len(texts) handled

RESULT SUMMARY

98.2% pass rate across 55 automated tests on live Render infrastructure.

# “The API is the product.”

The UI just calls it. The MCP server means any agent can run evals without touching a browser. That separation was harder to get right than any individual feature.



API      `criterion-api-c7mf.onrender.com/docs`

UI      `criterion-evall.streamlit.app`

WHY THIS MATTERS

The API-first design means any client, any agent, any interface can run evals. The MCP layer makes it native to agentic workflows with zero browser dependency.